

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import classification_report
from sklearn.metrics import plot_confusion_matrix
from sklearn.model_selection import train_test_split
from imblearn.over_sampling import RandomOverSampler
from imblearn.over_sampling import SMOTE
from sklearn import metrics
import plotly.offline as offline
import plotly.graph_objs as go
offline.init_notebook_mode()
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: data = pd.read_csv("TRAIN.csv")
```

```
In [3]: data.shape
```


```
Out[3]: (499120, 59)
```

```
In [4]: data.head()
```

```
Out[4]:
```

	id	target	ps_ind_01	ps_ind_02_cat	ps_ind_03	ps_ind_04_cat	ps_ind_05_cat	ps_ind_06_bin	ps_ind_0
0	7	0	2	2	5	1	0	0	
1	9	0	1	1	7	0	0	0	
2	13	0	5	4	9	1	0	0	
3	16	0	0	1	2	0	0	1	
4	17	0	0	2	0	1	0	1	

5 rows × 59 columns



```
In [5]: #check for null values
np.where(data.isnull())
```

```
Out[5]: (array([], dtype=int64), array([], dtype=int64))
```

```
In [6]: data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 499120 entries, 0 to 499119
Data columns (total 59 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                     499120 non-null  int64
1   target                 499120 non-null  int64
2   ps_ind_01              499120 non-null  int64
3   ps_ind_02_cat          499120 non-null  int64
4   ps_ind_03              499120 non-null  int64
5   ps_ind_04_cat          499120 non-null  int64
6   ps_ind_05_cat          499120 non-null  int64
7   ps_ind_06_bin          499120 non-null  int64
8   ps_ind_07_bin          499120 non-null  int64
9   ps_ind_08_bin          499120 non-null  int64
10  ps_ind_09_bin          499120 non-null  int64
11  ps_ind_10_bin          499120 non-null  int64
12  ps_ind_11_bin          499120 non-null  int64
13  ps_ind_12_bin          499120 non-null  int64
14  ps_ind_13_bin          499120 non-null  int64
15  ps_ind_14              499120 non-null  int64
16  ps_ind_15              499120 non-null  int64
17  ps_ind_16_bin          499120 non-null  int64
18  ps_ind_17_bin          499120 non-null  int64
19  ps_ind_18_bin          499120 non-null  int64
20  ps_reg_01              499120 non-null  float64
21  ps_reg_02              499120 non-null  float64
22  ps_reg_03              499120 non-null  float64
23  ps_car_01_cat          499120 non-null  int64
24  ps_car_02_cat          499120 non-null  int64
25  ps_car_03_cat          499120 non-null  int64
26  ps_car_04_cat          499120 non-null  int64
27  ps_car_05_cat          499120 non-null  int64
28  ps_car_06_cat          499120 non-null  int64
29  ps_car_07_cat          499120 non-null  int64
30  ps_car_08_cat          499120 non-null  int64
31  ps_car_09_cat          499120 non-null  int64
32  ps_car_10_cat          499120 non-null  int64
33  ps_car_11_cat          499120 non-null  int64
34  ps_car_11              499120 non-null  int64
35  ps_car_12              499120 non-null  float64
36  ps_car_13              499120 non-null  float64
37  ps_car_14              499120 non-null  float64
38  ps_car_15              499120 non-null  float64
39  ps_calc_01             499120 non-null  float64
40  ps_calc_02             499120 non-null  float64
41  ps_calc_03             499120 non-null  float64
42  ps_calc_04             499120 non-null  int64
43  ps_calc_05             499120 non-null  int64
44  ps_calc_06             499120 non-null  int64
45  ps_calc_07             499120 non-null  int64
46  ps_calc_08             499120 non-null  int64
47  ps_calc_09             499120 non-null  int64
48  ps_calc_10             499120 non-null  int64
49  ps_calc_11             499120 non-null  int64
50  ps_calc_12             499120 non-null  int64
51  ps_calc_13             499120 non-null  int64
52  ps_calc_14             499120 non-null  int64
53  ps_calc_15_bin         499120 non-null  int64

```

```

54 ps_calc_16_bin 499120 non-null int64
55 ps_calc_17_bin 499120 non-null int64
56 ps_calc_18_bin 499120 non-null int64
57 ps_calc_19_bin 499120 non-null int64
58 ps_calc_20_bin 499120 non-null int64
dtypes: float64(10), int64(49)
memory usage: 224.7 MB

```

```
In [7]: data.loc[data['target']==0].shape,data.loc[data['target']==1].shape
```


```
Out[7]: ((480895, 59), (18225, 59))
```

```
In [8]: data.describe()
```

```
Out[8]:
```

	id	target	ps_ind_01	ps_ind_02_cat	ps_ind_03	ps_ind_04_cat	ps_ir
count	4.991200e+05	499120.000000	499120.000000	499120.000000	499120.000000	499120.000000	499120.000000
mean	7.438299e+05	0.036514	1.901270	1.359304	4.421079	0.417220	0.417220
std	4.292680e+05	0.187566	1.983708	0.664709	2.699942	0.493389	0.493389
min	7.000000e+00	0.000000	0.000000	-1.000000	0.000000	-1.000000	-1.000000
25%	3.722780e+05	0.000000	0.000000	1.000000	2.000000	0.000000	0.000000
50%	7.436865e+05	0.000000	1.000000	1.000000	4.000000	0.000000	0.000000
75%	1.115324e+06	0.000000	3.000000	2.000000	6.000000	1.000000	1.000000
max	1.488027e+06	1.000000	7.000000	4.000000	11.000000	1.000000	1.000000

8 rows × 59 columns



```
In [9]: y = data['target']
X = data.drop('target',axis = 1)
```

```
In [10]: #oversampling on minority
smote = SMOTE()
X, y = smote.fit_resample(X, y)
```

```
In [11]: #train test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.1,stratify = y)
```

```
In [12]: X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
Out[12]: ((865611, 58), (96179, 58), (865611,), (96179,))
```

```
In [13]: X_train_id = X_train['id']
X_train = X_train.drop("id",axis=1)
```

```
X_test_id = X_test['id']
X_test = X_test.drop("id",axis=1)
```

Logistic Regression

```
In [14]: parameters = {
          'C': [1e-05, 0.0001, 0.001, 0.01, 0.1, 1, 10, 100, 1000, 10000]
        }
clf = LogisticRegression()
gscv = GridSearchCV(clf, param_grid = parameters, scoring='roc_auc', cv=10)
```

```
In [15]: gscv.fit(X_train,y_train)
```

```
Out[15]: GridSearchCV(cv=10, estimator=LogisticRegression(),
                    param_grid={'C': [1e-05, 0.0001, 0.001, 0.01, 0.1, 1, 10, 100,
                                      1000, 10000]},
                    scoring='roc_auc')
```

```
In [16]: result=pd.DataFrame(gscv.cv_results_)
         display(result)
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	params	split0_test_score	spl
0	9.009678	3.951964	0.056649	0.032646	0.00001	{'C': 1e-05}	0.899379	
1	7.278608	0.343338	0.045086	0.001564	0.0001	{'C': 0.0001}	0.946825	
2	7.542290	0.265399	0.045220	0.002442	0.001	{'C': 0.001}	0.955158	
3	7.769092	0.416858	0.047615	0.006377	0.01	{'C': 0.01}	0.958640	
4	7.796304	0.420956	0.048022	0.004345	0.1	{'C': 0.1}	0.960694	
5	7.960440	0.342552	0.049008	0.006694	1	{'C': 1}	0.958686	
6	7.930875	0.211714	0.048654	0.004804	10	{'C': 10}	0.958591	
7	7.991924	0.366423	0.046891	0.004362	100	{'C': 100}	0.960006	
8	8.023064	0.407536	0.047649	0.003574	1000	{'C': 1000}	0.958624	
9	7.544941	0.101527	0.045450	0.002535	10000	{'C': 10000}	0.958002	



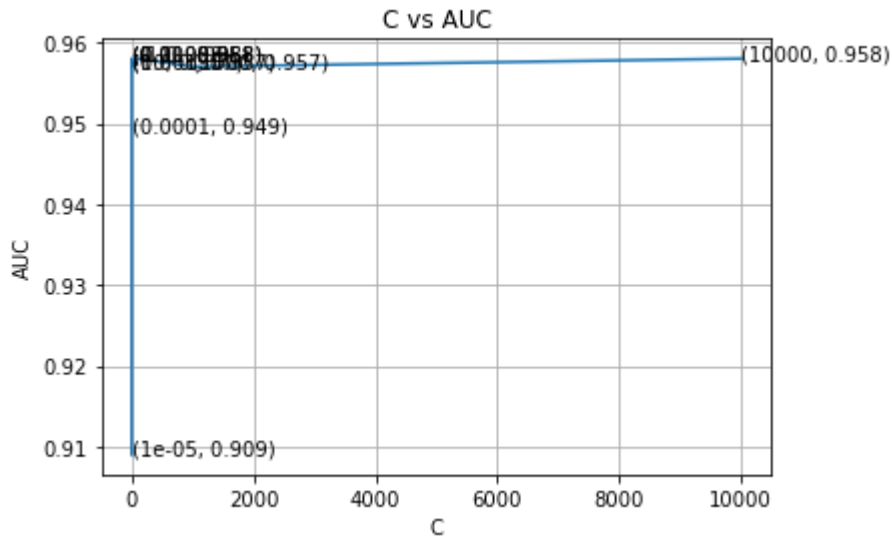
```
In [17]: C = result['param_C'].values
         avgaccscore = np.round(result['mean_test_score'].values,3)
         fig, ax = plt.subplots()
         sns.lineplot(x= C, y = avgaccscore)
         for i, txt in enumerate(avgaccscore):
```

```

ax.annotate((C[i],avgaccscore[i]), (C[i],avgaccscore[i]))
plt.grid()
plt.title("C vs AUC")
plt.xlabel("C")
plt.ylabel("AUC")
plt.show()

print("best hyperparameters :", gscv.best_params_)
print("Best AUC :",gscv.best_score_)

```



```

best hyperparameters : {'C': 1}
Best AUC : 0.9583216348124612

```

RandomForest Classifier

```

In [18]: parameters ={'max_depth': [1, 5],
                      'min_samples_split':[5, 10]}

clf = RandomForestClassifier()
gscv = GridSearchCV(clf, param_grid = parameters, scoring='roc_auc', cv=10)

```

```

In [19]: gscv.fit(X_train,y_train)

```

```

Out[19]: GridSearchCV(cv=10, estimator=RandomForestClassifier(),
                      param_grid={'max_depth': [1, 5], 'min_samples_split': [5, 10]},
                      scoring='roc_auc')

```


```

In [20]: result=pd.DataFrame(gscv.cv_results_)
display(result)

```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_max_depth	param_min_samples
0	17.864940	0.807276	0.314597	0.012424	1	
1	18.386464	0.993228	0.330494	0.028556	1	

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_max_depth	param_min_samples_
2	60.943874	0.548633	0.507395	0.008325	5	
3	62.731793	1.914924	0.528651	0.021751	5	



In [21]:


```
md = result['param_max_depth'].values
mss = result['param_min_samples_split'].values
avgaucscore = np.round(result['mean_test_score'].values,3)

trace1 = go.Scatter3d(x=md,y=mss,z=avgaucscore, name = 'Cross validation')

tc = [trace1]
layout = go.Layout(scene = dict(
    xaxis = dict(title='max_depth'),
    yaxis = dict(title='min_sample_split'),
    zaxis = dict(title='AUC'),))

fig = go.Figure(data=tc, layout=layout)
offline.iplot(fig, filename='3d-scatter-colorscale')

print("best hyperparameters :", gscv.best_params_)
print("Best AUC :",gscv.best_score_)
```



best hyperparameters : {'max_depth': 5, 'min_samples_split': 10}
Best AUC : 0.9623506480350908

Since Random Forest classifier is performing better than Logistic Regression, we can train the final model using best parameter from trained GridSearchCV on Random Forest Classifier.

```
In [22]: clf = RandomForestClassifier(max_depth = gscv.best_params_['max_depth'], min_samples_sp
         clf.fit(X_train,y_train)
```

```
Out[22]: RandomForestClassifier(max_depth=5, min_samples_split=10)
```

```
In [23]: #performance on test data
         y_test_pred = clf.predict_proba(X_test)[:,-1]
         fpr, tpr, thresholds = metrics.roc_curve(y_test.values, y_test_pred)
         auc_score = metrics.auc(fpr, tpr)
         print(auc_score)
```

0.9587476283890529

```
In [24]: score_data = pd.read_csv("SCORE.csv")
```

```
In [25]: score_data.head()
```

```
Out[25]:
```

	id	ps_ind_01	ps_ind_02_cat	ps_ind_03	ps_ind_04_cat	ps_ind_05_cat	ps_ind_06_bin	ps_ind_07_
0	1249566	0	1	9	0	0	1	
1	1249567	1	1	8	1	0	1	
2	1249568	0	1	7	0	4	0	
3	1249569	3	1	3	1	0	0	
4	1249570	4	1	10	0	0	0	

5 rows × 58 columns

```
In [26]: score_result = pd.DataFrame()
         score_result['id'] = score_data['id']
         score_data = score_data.drop('id',axis = 1)
```

```
In [27]: score_result['Probability'] = clf.predict_proba(score_data)[:,-1]
         score_result['PRED_Target'] = clf.predict(score_data)
```

```
In [28]: score_result.to_csv("result.csv")
```

